

Lab 1 – Create Azure resources

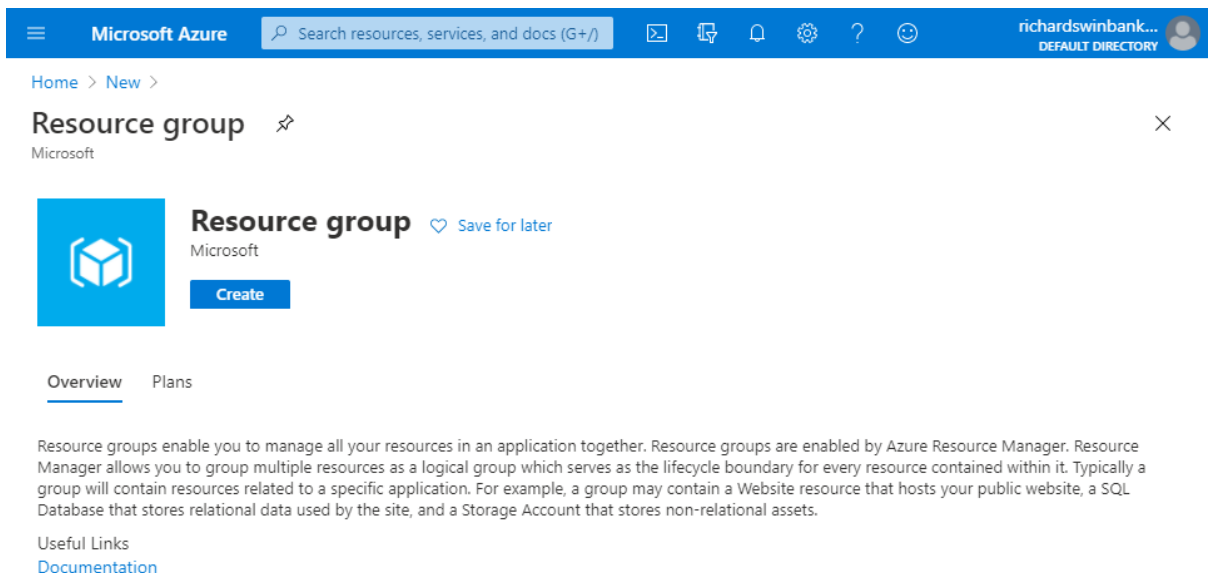
Welcome to Lab 1!

To complete the labs, you're going to need an Azure subscription (if you don't have one, you can sign up for a [free trial](#)). Whichever subscription you choose, you'll need enough access to create resources in it.

Lab 1.1 – Create a resource group

Resource groups are logical containers for resources in Azure. In this lab you will create a resource group to contain all the resources you create in later labs. This will make cleaning up easy – when you've finished all the labs, you can just delete the resource group.

1. In the [Azure Portal](#), click “Create a resource” and search for “Resource group”.
2. On the “Resource group” overview, click “Create”.




3. Give the resource group a name, and choose the Region geographically closest to you.
4. Click “Review + create”, then “Create”.

Lab 1.2 – Create data lake storage

Data lake storage is blob storage in an Azure storage account, with an important feature: hierarchical namespaces are **enabled**. This makes certain file operations – renaming file folders, for example – much more efficient.

1. In the portal, click “Create a resource” and search for “Storage account”. Click “Create” on the overview screen.
2. Complete the **Basics** tab like this:
 - Choose your subscription and the resource group you created in Lab 1.1.
 - Enter a storage account name – this must be globally unique (across the entire Azure platform).



- Choose the same location you specified for your resource group. Having storage located close to you makes for faster data transfers.
- Choose account kind “StorageV2” – this is the only kind that supports hierarchical namespaces.
- Choose replication option “Locally-redundant storage”. This is nice and cheap for lab work, but you’ll want something more resilient in a production environment!

[Home](#) > [New](#) > [Storage account](#) >

Create storage account

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription *

Resource group * [Create new](#)

Instance details

The default deployment model is Resource Manager, which supports the latest Azure features. You may choose to deploy using the classic deployment model instead. [Choose classic deployment model](#)

Storage account name * ✓

Location *

Performance ☒ Standard ☐ Premium

Account kind

Replication

Blob access tier (default) ☐ Cool ☒ Hot

[Review + create](#)

[< Previous](#)

[Next : Networking >](#)

3. On the **Advanced** tab (click through Networking and Data protection to get there), enable hierarchical namespaces. This step is **essential** to make this storage account a data lake.

[Basics](#) [Networking](#) [Data protection](#) [Advanced](#) [Tags](#) [Review + create](#)

Security

Secure transfer required ☐ Disabled ☒ Enabled

Allow Blob public access ☐ Disabled ☒ Enabled

Minimum TLS version

Infrastructure encryption ☒ Disabled ☐ Enabled

i Sign up is currently required to enable infrastructure encryption on a per-subscription basis. [Sign up for infrastructure encryption](#)

Azure Files


Large file shares ☒ Disabled ☐ Enabled

Data Lake Storage Gen2

Hierarchical namespace ☐ Disabled ☒ Enabled



NFS v3 ☒ Disabled ☐ Enabled

i Sign up is currently required to utilize the NFS v3 feature on a per-subscription basis. [Sign up for NFS v3](#)




- Click “Review + create”, then “Create”. When the data lake finishes deploying (this may take a couple of minutes), click on “Go to resource”.

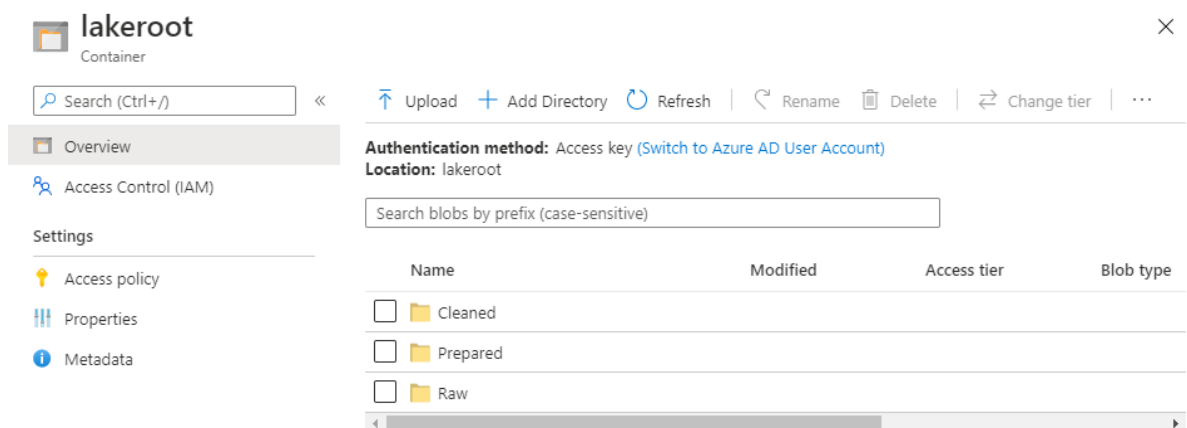
The storage account blade displays tiles for each of the four supported storage types. If your “Containers” tile looks like the one on the left, you’re in business. (If it looks like the one on the right, you forgot to enable hierarchical namespaces – delete the storage account and have another go).



- Click on the data lake containers tile to open the (empty) list of containers. Use the “+ Container” button to create a container with the name “lakeroot”. After creation, the container appears in the list.




- Click on the new “lakeroot” entry to open the container. The menu bar above the list now contains a “+ Add Directory” button – use this to create three directories in the container: “Raw”, “Cleaned” and “Prepared”.



Lab 1.3 – Create an Azure Data Factory

The main event! It’s time to create your Azure Data Factory instance.

- In the portal, click “Create a resource” and search for “Data Factory”. Click “Create” on the overview screen.



2. Complete the **Basics** tab like this:

- Choose your subscription and the resource group you created in Lab 1.1.
- Choose the same location (region) you specified for your storage account. **This has cost implications** – transferring data from a storage account in one region to a data factory in another incurs an outbound data transfer charge.
- Enter a data factory name – this must be globally unique.
- Choose version V2.

Home > New > Data Factory >

Create Data Factory

Basics Git configuration Tags Review + create

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * ⓘ Free Trial ▼

Resource group * ⓘ a-day-full-of-adf ▼
[Create new](#)

Instance details

Region * ⓘ UK South ▼


Name * adayfullof-adf ✓

Version * ⓘ V2 ▼

Enable Managed Virtual Network (Preview) ⓘ ☐

[Review + create](#) < Previous Next : Git configuration >


3. If you have a Git repository available, you can connect your new factory to it on the **Git configuration** tab. Don't worry if you don't have a repo ready – just click "Configure Git later".
4. Click "Review + create", then "Create". When the factory deployment is complete, click on "Go to resource", then on the "Author & Monitor" tile to launch the ADF User Experience (ADF UX).

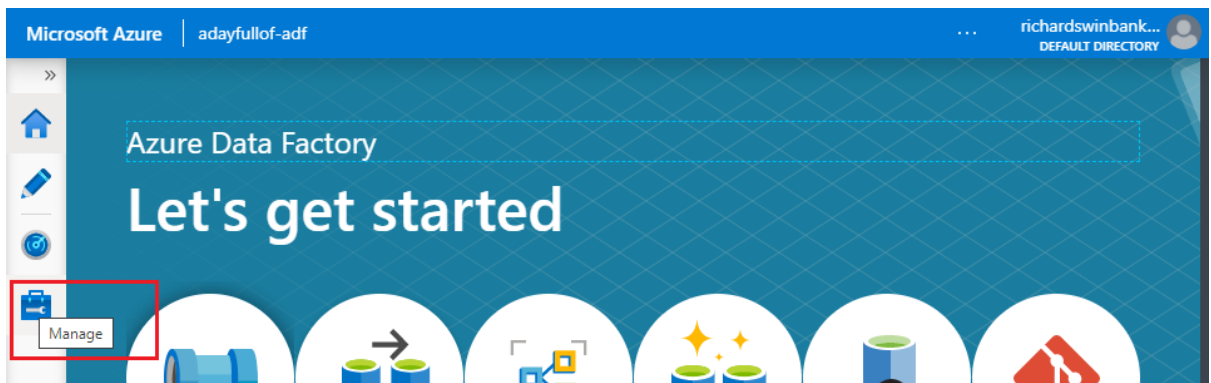


Lab 1.4 – Connect to the data lake

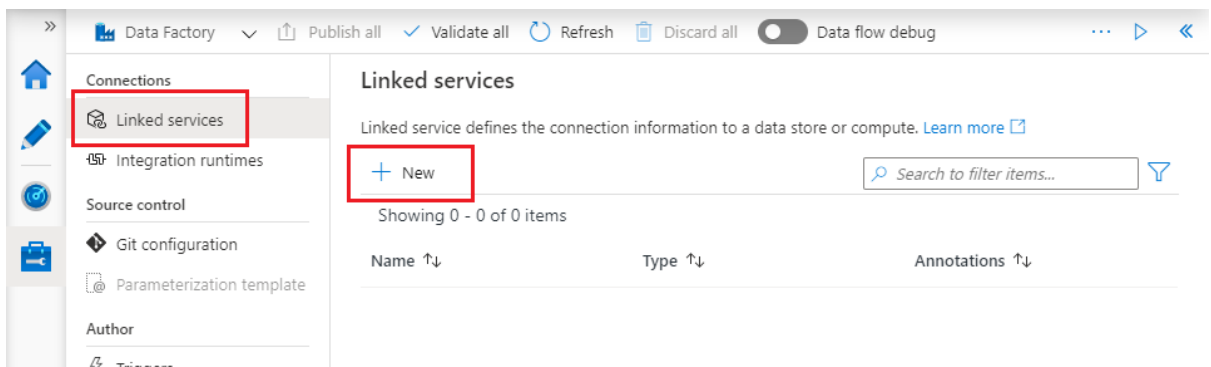
To enable ADF pipelines to use data in the lake you will need a data factory **linked service** connection.

1. In the ADF UX's leftmost sidebar, click the "Manage" button to open the Management Hub.

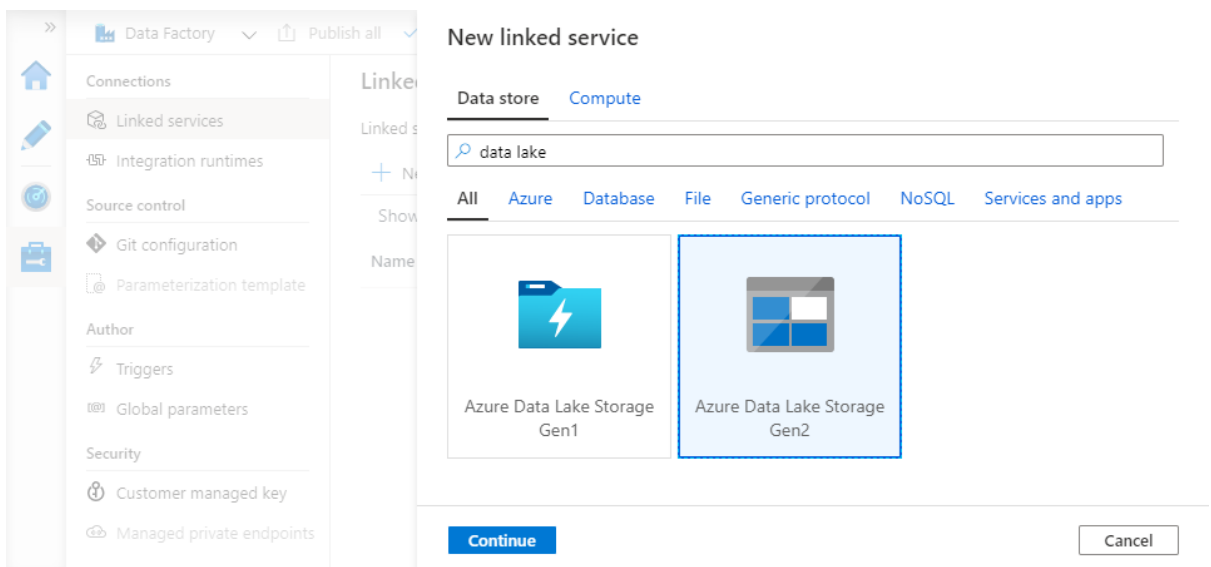




2. Select the “Linked services” item from the “Connections” section of the Management Hub sidebar, then in the main Linked services pane click “+ New”.



3. Search for “data lake”, then choose “Azure Data Lake Storage Gen2” and click “Continue”.



4. Configure linked service details on the “New linked service (Azure Data Lake Storage Gen2)” blade like this:

- Give it a name.
- Choose Authentication method “Managed Identity”. The default method (“Account key”) requires extra work to pass keys around securely – “Managed Identity” uses an Azure Active Directory service principal created automatically for your data factory when you created it. This service principal is called a Managed Service Identity (MSI).
- Choose your data lake storage account from the “Storage account name” dropdown.

- Click “Test connection” at the bottom of the blade. The connection test will fail, because the factory’s MSI does not have access to the data lake yet – you will receive an error message like “This request is not authorized to perform this operation using this permission”.


The screenshot shows the Azure Data Factory portal interface. On the left is a navigation pane with options like Home, Connections, Linked services, Integration runtimes, Source control, Git configuration, Parameterization template, Author, Triggers, Global parameters, Security, Customer managed key, and Managed private endpoints. The main area displays the 'New linked service (Azure Data Lake Storage Gen2)' blade. It includes fields for 'Connect via integration runtime' (set to 'AutoResolveIntegrationRuntime'), 'Authentication method' (set to 'Managed Identity'), and 'Account selection method' (set to 'From Azure subscription'). Below these are fields for 'Azure subscription' (set to 'Select all') and 'Storage account name' (set to 'adayfullof-adf'). The 'Managed identity name' is 'adayfullof-adf' and the 'Managed identity object ID' is 'c193f56d-6...'. The 'Test connection' button is highlighted with a red box. An error message is displayed, stating: 'ADLS Gen2 operation failed for: Operation returned an invalid status code 'Forbidden'. Account: 'adayfullof-adf', Filesystem: 'filesystem'. ErrorCode: 'AuthorizationPermissionMismatch'. Message: 'This request is not authorized to perform this operation using this permission.'. RequestId: 'bce1ce45-601f-00b-3ca1-80594b000000'. TimeStamp: 'Tue, 01 Sep 2020 20:50:28 GMT'. Operation returned an invalid status code 'Forbidden' Activity ID: 24dff969-fbfa-42b5-b154-f7a54615559d.' The error message is also highlighted with a red box. The 'Create' button is visible at the bottom of the blade.

- You will grant the necessary access in a moment – for now, just click “Create” to create the linked service.


Lab 1.5 – Grant access to the data lake

You can manage access to Azure resources in the Azure portal. Open a new browser tab to allow you to keep the ADF UX open.

- Browse to your data lake resource blade – you can find it in the list of resources on the portal home page, or by selecting your resource group to see resources inside it, or by using the search box in the portal’s top menu bar.
- Click “Access control (IAM)” in the storage account (data lake) resource blade.
- In the “Add a role assignment” tile, click the “Add” button.
- In the “Add role assignment” blade, choose the “Storage Blob Data Contributor” role – this authorises read, write and delete access in your data lake.




- Enter the name of your data factory in the “Select” text box to search for the MSI. (The MSI’s display name is the same as the data factory name). When the MSI appears below the “Select” box, click to select it, then click “Save”.



The screenshot shows the Azure portal interface for a storage account named 'adayfullofadsfa'. The 'Access control (IAM)' section is active. A modal window titled 'Add role assignment' is open. In this window, the 'Role' is set to 'Storage Blob Data Contributor'. The 'Assign access to' dropdown is set to 'Azure AD user, group, or service principal'. The 'Select' dropdown shows a search for 'adayfullof', and the result 'adayfullof-adf' is highlighted with a red rectangular box. Below the search results, it says 'Selected members: No members selected. Search for and add one or more members you want to assign to the role for this resource.' At the bottom of the modal, there are 'Save' and 'Discard' buttons.

- Return to the ADF UX Management Hub, and click on your data lake linked service to re-open the editing blade. Click “Test connection” again, and verify that this time the connection test succeeds. If the test fails, wait a few minutes, then try again – it may take a short time for permission changes to take effect). Click “Cancel” to close the editing blade.
- If you have not enabled Git for your data factory, the only way to save your changes is to publish them – click “Publish all” in the top menu bar to do so. (If Git is enabled, linked service changes are saved automatically to your Git repo).



The screenshot shows the Azure portal interface for a data factory. The top menu bar includes buttons for 'Data Factory', 'Publish all' (highlighted with a red box), 'Validate all', 'Refresh', 'Discard all', and 'Data flow debug'. The left sidebar shows the 'Connections' section with 'Linked services' selected. The main area displays the 'Linked services' section, which includes a 'New' button, a search filter, and a list of linked services. The list shows one item: 'AzureDataLak...' with a type of 'Azure Data Lake Storage Gen2'.

Recap

In Lab 1 you:

- created an Azure resource group
- created data lake storage
- created an instance Azure Data Factory
- created and authorised a connection from the factory to your data lake storage.

